

★ Your membership will expire on August 4, 2024 [Reactivate membership](#)

★ Member-only story

# iTransformer: The Latest Breakthrough in Time Series Forecasting


Discover the architecture of iTransformer and apply the model in a small experiment using Python.



Marco Peixeiro  · Follow

Published in Towards Data Science

9 min read · Apr 9, 2024



Listen



Share



More



Photo by [David Clode](#) on [Unsplash](#)

The field of forecasting has seen a lot of activity in the realm of foundation models, with models like [Lag-LLaMA](#), [Time-LLM](#), [Chronos](#) and Moirai being proposed since the beginning of 2024.

However, their performance has been a bit underwhelming (for reproducible benchmarks, see [here](#)), and I believe that data-specific models are still the optimal solution at the moment.

To that end, the Transformer architecture has been applied in many forms for time series forecasting, with [PatchTST](#) achieving state-of-the-art performance for long-horizon forecasting.

Challenging PatchTST comes the **iTransformer** model, proposed in March 2024 in the paper [iTransformer: Inverted Transformers Are Effective for Time Series Forecasting](#).

In this article, we discover the strikingly simple concept behind iTransformer and explore its architecture. Then, we apply the model in a small experiment and compare its performance to [TSMixer](#), [N-HiTS](#) and PatchTST.

For more details, make sure to read the [original paper](#).

Learn the latest time series analysis techniques with my [free time series cheat sheet](#) in Python! Get the implementation of statistical and deep learning techniques, all in Python and TensorFlow!


Let's get started!

## Explore iTransformer

The idea behind iTransformer comes from the realization that the vanilla Transformer model uses temporal *tokens*.

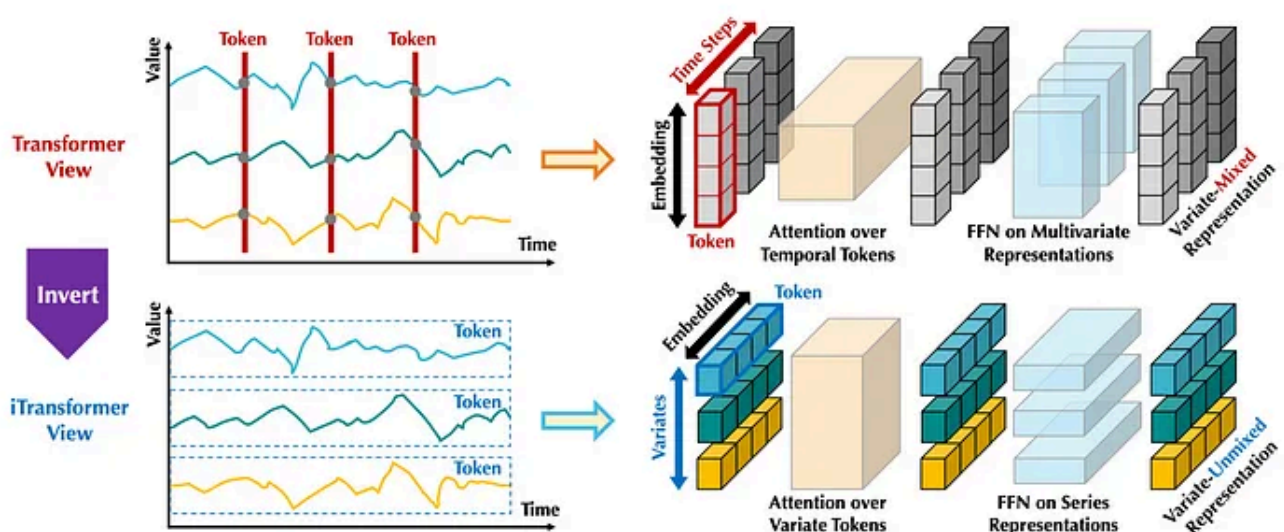
This means that the model looks at all features at a single time step. Thus, it is challenging for the model to learn temporal dependencies when looking at one time step at a time.

A solution to that problem is patching, which was proposed with the PatchTST model. With patching, we simply group time points together before tokenizing and embedding them, as shown below.



Visualizing patching. Here, we have a sequence of 15 timesteps, with a patch length of 5 and a stride of 5 as well, resulting in three patches. Image by the author.

In iTransformer, we push patching to the extreme by simply applying the model on the inverted dimensions.



The iTransformer is applied on the inverted shape of the series. That way, the whole feature across time is tokenized and embedded. Image by Y. Liu, T. Hu, H. Zhang, H. Wu, S. Wang, L. Ma, M. Long from [iTransformer: Inverted Transformers Are Effective for Time Series Forecasting](https://arxiv.org/abs/2303.05499).

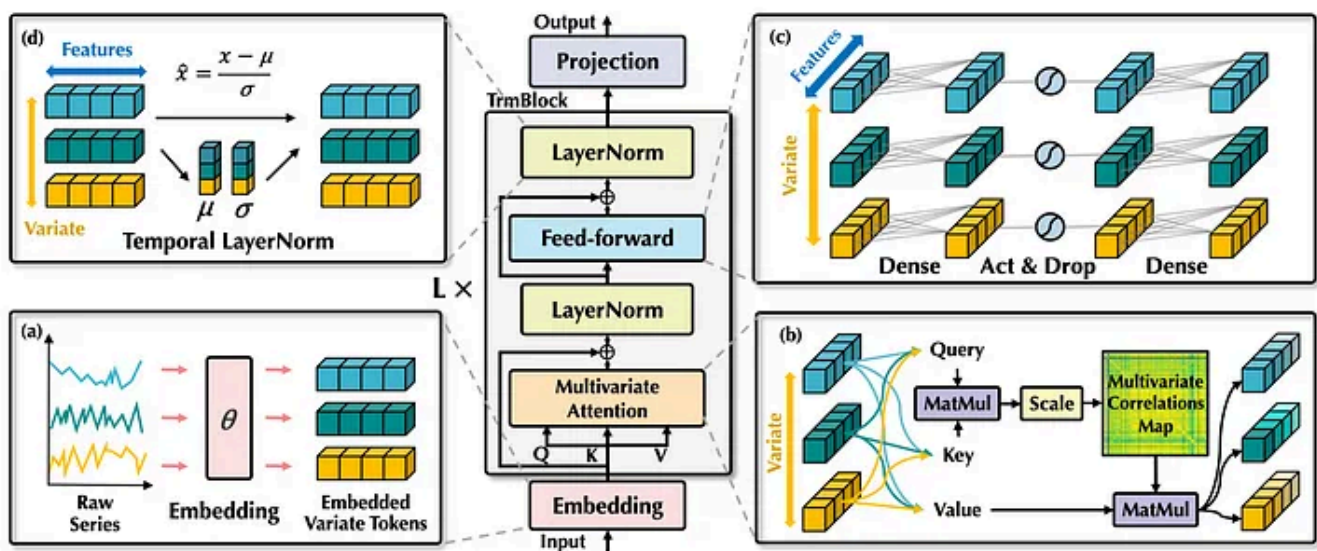
In the figure above, we can see how the iTransformer differs from the vanilla Transformer. Instead of looking at all features at one time step, it looks at one feature across many time steps. This is done simply by inverting the shape of the input.

This way, the attention layer can learn multivariate correlations and the feed-forward network encodes the representation of the whole input sequence.

Now that we grasp the general idea behind iTransformer, let's take a look at its architecture in more detail.

### Architecture of iTransformer

The iTransformer employs the vanilla encoder-decoder architecture with the embedding, projection and Transformer blocks, as originally proposed in the seminal paper Attention Is All You Need in 2017.



Architecture of iTransformer. Image by Y. Liu, T. Hu, H. Zhang, H. Wu, S. Wang, L. Ma, M. Long from iTransformer: Inverted Transformers Are Effective for Time Series Forecasting.

From the figure above, we can see that the building blocks are the same, but their function is entirely different. Let's take a closer look.

### Embedding layer

First, the input series are independently embedded as tokens. Again, this is like an extreme case of patching, where instead of tokenizing subsequences of the input, the model tokenizes the entire input sequence.

### Multivariate attention

Then, the embeddings are sent to the attention layer, where it will learn a multivariate correlation map.

This is possible because the inverted model sees each feature as an independent process. As a result, the attention mechanism learns correlations between pairs of features, making iTransformer especially suitable for multivariate forecasting tasks.

### **Layer normalization**

The output of the attention layer is sent to the normalization layer.

In the traditional Transformer architecture, normalization is done for all features at a fixed timestamp. This can introduce interaction noise, meaning that the model is learning useless relationships. Plus, it can result in an overly smooth signal.

By contrast, since iTransformer inverts the dimensions, normalization is done across timestamps. This helps the model tackle non-stationary series, and it reduces the noise in the series.

### **Feed-forward network**

Finally, the feed-forward network (FFN) learn a deep representation of the incoming tokens.

Again, since the shapes are inverted, the multilayer perceptron (MLP) can learn different temporal properties, like periodicity and amplitude. This mimics the capabilities of MLP-based models, like N-BEATS, N-HiTS and TSMixer.

### **Projection**

From here, it is simply a matter of stacking many blocks composed of:

- attention layer
- layer normalization
- feed-forward network
- layer normalization

Each block learns a different representation of the input series. Then, the output of the stack of blocks is sent through a linear projection step to obtain the final

forecasts.

In summary, the iTransformer is not a new architecture; it does not reinvent the Transformer. It simply applies it on the inverted dimensions of the input, which allows the model to learn multivariate correlations and capture temporal properties.

Now that we have a deep understanding of the iTransformer model, let's apply it in a small forecasting experiment.

## Forecasting with iTransformer

For this small experiment, we apply the iTransformer model on [Electricity Transformer dataset](#) released under the Creative Commons License.

This is a popular benchmark dataset that tracks the oil temperature of an electricity transformer from two regions in a province of China. For both regions, we have a dataset sampled at each hour and every 15 minutes, for a total of four datasets.

While the iTransformer is fundamentally a multivariate model, we test its univariate forecasting capabilities on a horizon of 96 time steps.

The code for this experiment is available on [GitHub](#).

Let's get started!

### Initial setup

For this experiment, we use the library *neuralforecast*, as I believe it provides the fastest and most intuitive ready-to-use implementations of deep learning methods.

```
import pandas as pd
import numpy as np
import matplotlib.pyplot as plt

from datasetsforecast.long_horizon import LongHorizon

from neuralforecast.core import NeuralForecast
from neuralforecast.models import NHITS, PatchTST, iTransformer, TSMixer
```

Note that at the time of writing this article, iTransformer is not available in a public release of *neuralforecast* just yet. To access the model immediately you can run:

```
pip install git+https://github.com/Nixtla/neuralforecast.git
```

Now, let's write a function to load the ETT datasets, along with their validation size, test size, and frequency.

[Open in app](#)**Medium** Search

```
Y_df = Y_df[Y_df['unique_id'] == 'OT']
Y_df['ds'] = pd.to_datetime(Y_df['ds'])
val_size = 11520
test_size = 11520
freq = '15T'
elif name == "ettm2":
    Y_df, *_ = LongHorizon.load(directory='.', group='ETTM2')
    Y_df = Y_df[Y_df['unique_id'] == 'OT']
    Y_df['ds'] = pd.to_datetime(Y_df['ds'])
    val_size = 11520
    test_size = 11520
    freq = '15T'
elif name == 'etth1':
    Y_df, *_ = LongHorizon.load(directory='.', group='ETTh1')
    Y_df['ds'] = pd.to_datetime(Y_df['ds'])
    val_size = 2880
    test_size = 2880
    freq = 'H'
elif name == "etth2":
    Y_df, *_ = LongHorizon.load(directory='.', group='ETTh2')
    Y_df['ds'] = pd.to_datetime(Y_df['ds'])
    val_size = 2880
    test_size = 2880
    freq = 'H'

return Y_df, val_size, test_size, freq
```

The function above conveniently loads the data in the format expected by *neuralforecast*, where we have a `unique_id` column to label unique time series, a `ds` column for the timestamp, and a `y` column with the value of our series.

Also note that the validation and test sizes are consistent with what the scientific community uses when publishing their papers.



We are now ready to train the models.

### Training and forecasting

To train the iTransformer model, we simply need to specify the:

- forecast horizon
- input size
- number of series

Remember that iTransformer is fundamentally a multivariate model, which is why we need to specify the number of series when fitting the model.

Since we are in a univariate scenario, `n_series=1`.

```
iTransformer(h=horizon,
             input_size=3*horizon,
             n_series=1,
             max_steps=1000,
             early_stop_patience_steps=3)
```

In the code block above, we also specify the maximum number of training steps, and we set early stopping to three iterations to avoid overfitting.

We then do the same for the other models, and place them in a list.

```
horizon = 96

models = [
    iTransformer(h=horizon, input_size=3*horizon, n_series=1, max_steps=1000, e
    TSMixer(h=horizon, input_size=3*horizon, n_series=1, max_steps=1000, early_
    NHITS(h=horizon, input_size=3*horizon, max_steps=1000, early_stop_patience_
    PatchTST(h=horizon, input_size=3*horizon, max_steps=1000, early_stop_patier
]
```

Great! Now, we simply initialize the `NeuralForecast` object which gives access to methods for training, cross-validation, and making predictions.



```

nf = NeuralForecast(models=models, freq=freq)
nf_preds = nf.cross_validation(df=Y_df, val_size=val_size, test_size=test_size,

```

Finally, we evaluate the performance of each model using the `utilsforecast` library.

```

from utilsforecast.losses import mae, mse
from utilsforecast.evaluation import evaluate

ettm1_evaluation = evaluate(df=nf_preds, metrics=[mae, mse], models=['iTransformer'])
ettm1_evaluation.to_csv('ettm1_results.csv', index=False, header=True)

```

These steps are then repeated for all datasets. The complete function to run this experiment is shown below.

```

from utilsforecast.losses import mae, mse
from utilsforecast.evaluation import evaluate

datasets = ['ettm1', 'ettm2', 'etth1', 'etth2']

for dataset in datasets:

    Y_df, val_size, test_size, freq = load_data(dataset)

    horizon = 96

    models = [
        iTransformer(h=horizon, input_size=3*horizon, n_series=1, max_steps=1000, early_stop_patience=10),
        TSMixer(h=horizon, input_size=3*horizon, n_series=1, max_steps=1000, early_stop_patience=10),
        NHITS(h=horizon, input_size=3*horizon, max_steps=1000, early_stop_patience=10),
        PatchTST(h=horizon, input_size=3*horizon, max_steps=1000, early_stop_patience=10)
    ]

    nf = NeuralForecast(models=models, freq=freq)
    nf_preds = nf.cross_validation(df=Y_df, val_size=val_size, test_size=test_size)
    nf_preds = nf_preds.reset_index()

    evaluation = evaluate(df=nf_preds, metrics=[mae, mse], models=['iTransformer'])
    evaluation.to_csv(f'{dataset}_results.csv', index=False, header=True)

```

Once this is done running, we have predictions from all of our models on all datasets. We can then move on to evaluation.

## Performance evaluation

Since we saved all performance metrics in CSV files, we can read them using *pandas* and plot the performance of each model for each dataset.

```
files = ['etth1_results.csv', 'etth2_results.csv', 'ettm1_results.csv', 'ettm2_
datasets = ['etth1', 'etth2', 'ettm1', 'ettm2']

dataframes = []

for file, dataset in zip(files, datasets):
    df = pd.read_csv(file)
    df['dataset'] = dataset

    dataframes.append(df)

full_df = pd.concat(dataframes, ignore_index=True)
full_df = full_df.drop(['unique_id'], axis=1)
```

Then, to plot the metrics:

```
import matplotlib.pyplot as plt
import numpy as np

dataset_names = full_df['dataset'].unique()
model_names = ['iTransformer', 'TSMixer', 'NHITS', 'PatchTST']

fig, axs = plt.subplots(2, 2, figsize=(15, 15))
bar_width = 0.35

axs = axs.flatten()

for i, dataset_name in enumerate(dataset_names):
    df_subset = full_df[(full_df['dataset'] == dataset_name) & (full_df['metric'] == 'mae')]
    mae_vals = df_subset[model_names].values.flatten()
    df_subset = full_df[(full_df['dataset'] == dataset_name) & (full_df['metric'] == 'mse')]
    mse_vals = df_subset[model_names].values.flatten()

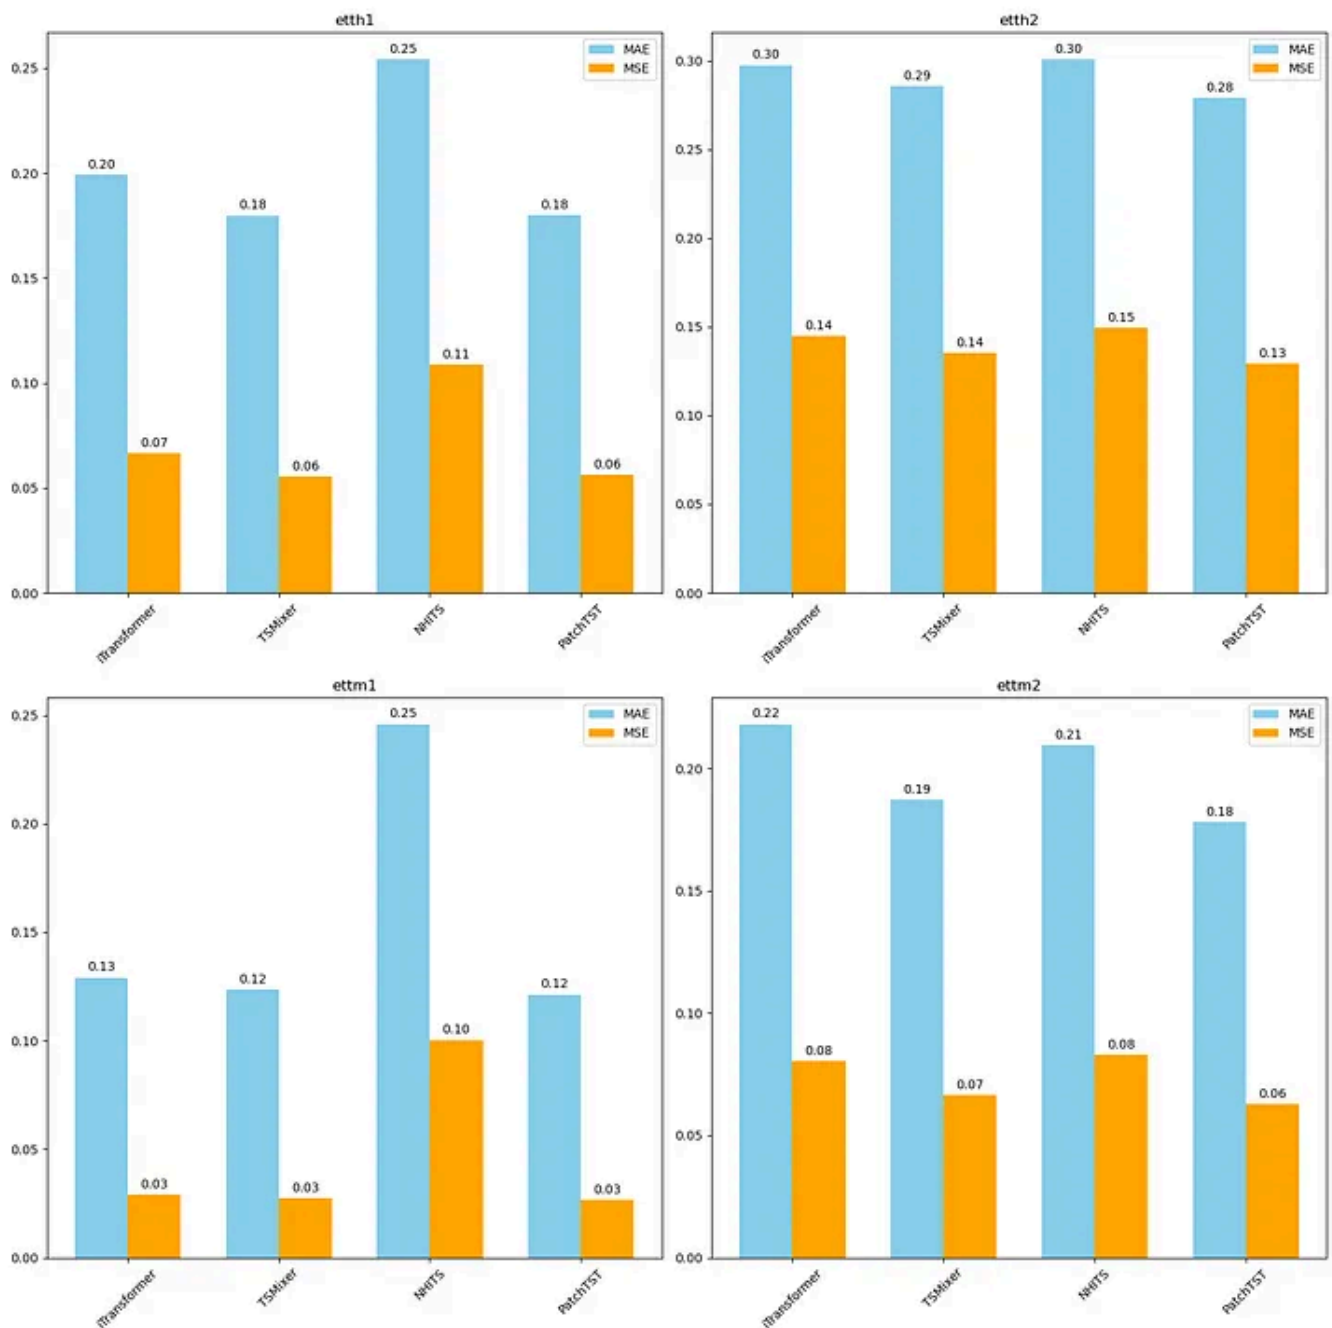
    indices = np.arange(len(model_names))

    bars_mae = axs[i].bar(indices - bar_width / 2, mae_vals, bar_width, color='red')
    bars_mse = axs[i].bar(indices + bar_width / 2, mse_vals, bar_width, color='blue')
```

```
for bars in [bars_mae, bars_mse]:
    for bar in bars:
        height = bar.get_height()
        axs[i].annotate(f'{height:.2f}',
                        xy=(bar.get_x() + bar.get_width() / 2, height),
                        xytext=(0, 3),
                        textcoords="offset points",
                        ha='center', va='bottom')

    axs[i].set_xticks(indices)
    axs[i].set_xticklabels(model_names, rotation=45)
    axs[i].set_title(dataset_name)
    axs[i].legend(loc='best')

plt.tight_layout()
```



MAE and MSE from all models on the ETT dataset for a horizon of 96 time steps. Image by the author.

From the figure above, we can see that the iTransformer performs fairly well on all datasets, but TSMixer is overall slightly better than iTransformer, and PatchTST is the overall champion model in this experiment.

Of course, keep in mind that we did not leverage the multivariate capabilities of iTransformer, and we only tested on a single forecast horizon. Therefore, it is not a complete assessment of the iTransformer's performance.

Nevertheless, it is interesting to see the model perform very closely to PatchTST, which further supports the idea that grouping time steps together before tokenizing

unlocks new heights in performance when using Transformers for time series forecasting.

## Conclusion

The iTransformer takes the vanilla Transformer architecture and simply applies it to the inverted shape of the input series.

That way, the entire series is tokenized, which mimics an extreme case of patching, as proposed in PatchTST.

This allows the model to learn multivariate correlations using its attention mechanism, while the feed-forward network learn temporal properties of the series.

The iTransformer was shown to achieve state-of-the-art performance for long-horizon forecasting on many benchmark datasets, although in our limited experiment, PatchTST performed best overall.

I firmly believe that every problem requires its unique solution, and now you can add the iTransformer to your toolbox and apply it in your projects.

Thanks for reading! I hope that you enjoyed it and that you learned something new!

---

*Looking to master time series forecasting? The check out [Applied Time Series Forecasting in Python](#). This is the only course that uses Python to implement statistical, deep learning and state-of-the-art models in 15 guided hands-on projects.*

---

Cheers 🍷

## Support me

Enjoying my work? Show your support with [Buy me a coffee](#), a simple way for you to encourage me, and I get to enjoy a cup of coffee! If you feel like it, just click the button below 🙌



## References

# iTransformer: Inverted Transformers Are Effective for Time Series Forecasting by Yong Liu, Tengge Hu, Haoran Zhang, Haixu Wu, Shiyu Wang, Lintao Ma, Mingsheng Long

Forecasting

Data Science

Machine Learning

Artificial Intelligence

Python



Follow

**Written by Marco Peixeiro**

22K Followers · Writer for Towards Data Science

Senior data scientist | Author | Instructor. I write hands-on articles with a focus on practical skills.

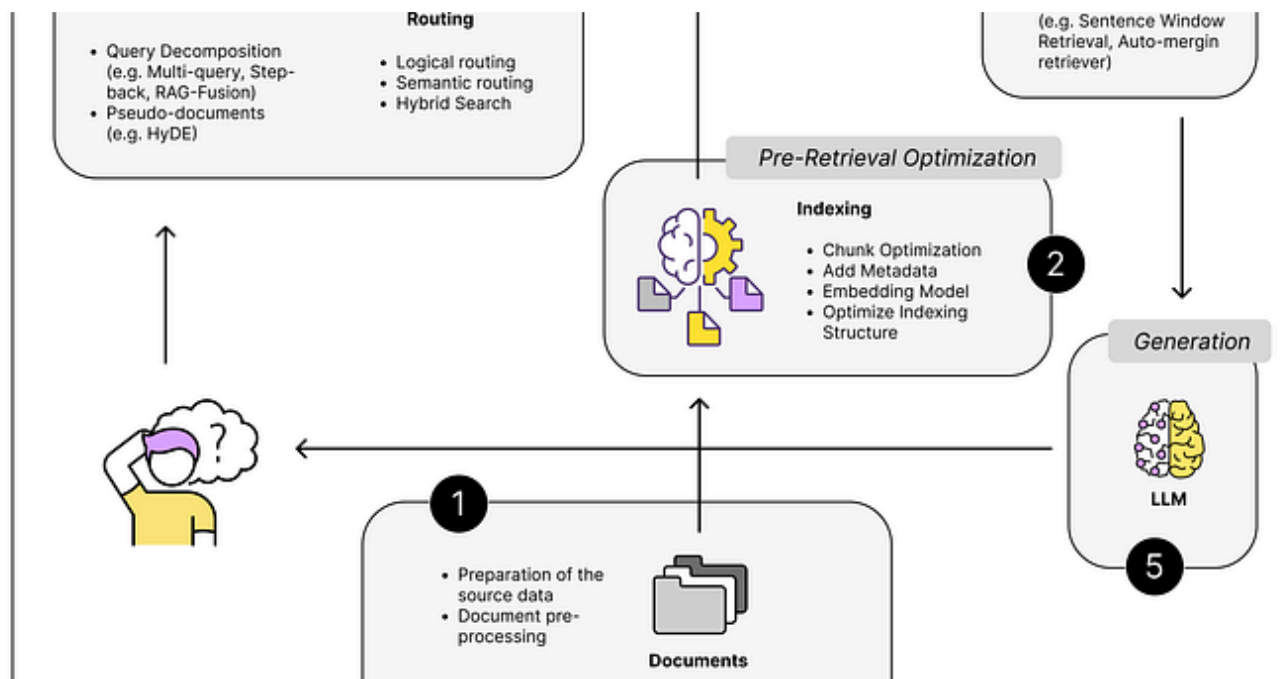
**More from Marco Peixeiro and Towards Data Science**

Marco Peixeiro in Towards Data Science

## SOFTS: The Latest Innovation in Time Series Forecasting

Discover the architecture of SOFTS and the novel STAD module, and apply it in a forecasting project using Python.

★ Jun 11 🤝 337



Dominik Polzer in Towards Data Science

## 17 (Advanced) RAG Techniques to Turn Your LLM App Prototype into a Production-Ready Solution

A collection of RAG techniques to help you develop your RAG app into something robust that will last

★ Jun 26 🤝 1.7K 💬 16







Almog Baku in Towards Data Science

## Building LLM Apps: A Clear Step-By-Step Guide

Comprehensive Steps for Building LLM-Native Apps: From Initial Idea to Experimentation, Evaluation, and Productization

Jun 10 🖱️ 6K 💬 17



Marco Peixeiro  in Towards Data Science

## Kolmogorov-Arnold Networks (KANs) for Time Series Forecasting

Discover the Kolmogorov-Arnold Networks (KANs) and apply them for time series forecasting using Python

★ May 14

👤 764

💬 6

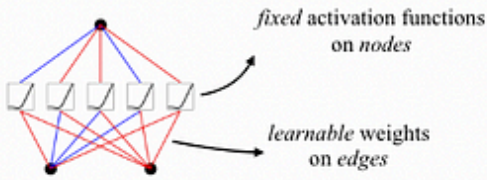
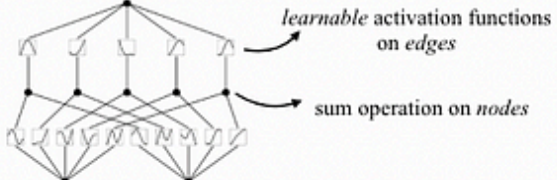
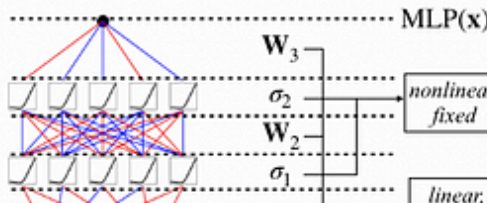
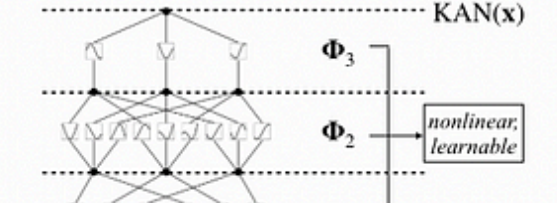
🔖

⋮

See all from Marco Peixeiro

See all from Towards Data Science

Recommended from Medium

Theorem	Universal Approximation Theorem	Kolmogorov-Arnold Representation Theorem
Formula (Shallow)	$f(\mathbf{x}) \approx \sum_{i=1}^{N(c)} a_i \sigma(\mathbf{w}_i \cdot \mathbf{x} + b_i)$	$f(\mathbf{x}) = \sum_{q=1}^{2n+1} \Phi_q \left( \sum_{p=1}^n \phi_{q,p}(x_p) \right)$
Model (Shallow)	(a)  fixed activation functions on nodes learnable weights on edges	(b)  learnable activation functions on edges sum operation on nodes
Formula (Deep)	$\text{MLP}(\mathbf{x}) = (\mathbf{W}_3 \circ \sigma_2 \circ \mathbf{W}_2 \circ \sigma_1 \circ \mathbf{W}_1)(\mathbf{x})$	$\text{KAN}(\mathbf{x}) = (\Phi_3 \circ \Phi_2 \circ \Phi_1)(\mathbf{x})$
Model (Deep)	(c)  MLP(x) W3, sigma2, W2, sigma1 nonlinear, fixed linear	(d)  KAN(x) Phi3, Phi2 nonlinear, learnable

 Theo Wolf in Towards Data Science

Kolmogorov-Arnold Networks: the latest advance in Neural Networks, simply explained

The new type of network that is making waves in the ML world.

★ May 13


👤 2.2K

💬 19

🔖

⋮





Shenggang Li in DataDrivenInvestor

## Time Series Forecasting: A Comparative Analysis of SARIMAX, RNN, LSTM, Prophet, and Transformer...

Assessing the Efficiency and Efficacy of Leading Forecasting Algorithms Across Diverse Datasets



Apr 13




409



4



### Lists




#### Predictive Modeling w/ Python

20 stories · 1366 saves



#### Practical Guides to Machine Learning

10 stories · 1650 saves




#### Natural Language Processing

1568 stories · 1120 saves



#### ChatGPT prompts

48 stories · 1777 saves



# A JOURNAL NOTEBOOK

## YOUR MODEL IN



Mandar Karhade, MD. PhD. in Towards AI

## Why RAG Applications Fail in Production

It worked as a prototype; then all went down!



Mar 20



2K



22



**Amazon.com**

*Software Development Engineer*

Seattle, WA

Mar. 2020 – May 2021

- Developed Amazon checkout and payment services to handle traffic of 10 Million daily global transactions
- Integrated Iframes for credit cards and bank accounts to secure 80% of all consumer traffic and prevent CSRF, cross-site scripting, and cookie-jacking
- Led Your Transactions implementation for JavaScript front-end framework to showcase consumer transactions and reduce call center costs by \$25 Million
- Recovered Saudi Arabia checkout failure impacting 4000+ customers due to incorrect GET form redirection

### Projects

**NinjaPrep.io (React)**

- Platform to offer coding problem practice with built in code editor and written + video solutions in React
- Utilized Nginx to reverse proxy IP address on Digital Ocean hosts
- Developed using Styled-Components for 95% CSS styling to ensure proper CSS scoping
- Implemented Docker with Seccomp to safely run user submitted code with < 2.2s runtime

**HeatMap (JavaScript)**

- Visualized Google Takeout location data of location history using Google Maps API and Google Maps heatmap code with React
- Included local file system storage to reliably handle 5mb of location history data
- Implemented Express to include routing between pages and jQuery to parse Google Map and implement heatmap overlay



Alexander Nguyen in Level Up Coding

## The resume that got a software engineer a \$300,000 job at Google.

1-page. Well-formatted.



Jun 1



12.4K



165



```
*_, a, b, *_ = [1, 2, 3, 4, 5, 6]
print(__, _)
```

What does this print?

- A) Syntax error
- B) [1] [4, 5, 6]
- C) [1, 2] [5, 6]
- D) [1, 2, 3] [6]
- E) <generator object <genexpr> at 0x1003847c0>

 Liu Zuo Lin

## You're Decent At Python If You Can Answer These 7 Questions Correctly

# No cheating pls!!

★ Mar 6 🤝 6K 💬 29



 Dr. Ernesto Lee

## Advanced Stock Pattern Prediction using LSTM with the Attention Mechanism in TensorFlow: A step by...

Introduction

Apr 8

 1K

 21





See more recommendations